



WiFi Password:  
W3lk0m@ArenbergB1



# Postgres on Kubernetes Workshop Leuven

6 June | Leuven, Belgium

Borys Neselovskyi  
Senior Sales Engineer, EDB

| Piotr Kolodziej  
Sales Engineer, EDB

Wifi netwerk:  
B1 Meeting Room

Wachtwoord:  
W3lkom@ArenbergB1



# Agenda

Start	End	Session
09:00	09:30	Registration & Welcome
09:30	09:45	Piros and Zebanza Introduction
09:45	10:00	Red Hat OpenShift & EDB Partnership
10:00	10:15	Introduction to CloudNativePG and EDB
10:15	10:45	CNPG Operator Reference Architecture
10:45	11:15	Functionalities for CNPG
11:15	13:00	Interactive session & demo
13:00	14:00	Lunch



# Piros & Zebanza Introduction





Premier

**Business Partner**  
Solution Provider

Datacenter infrastructure  
Middleware solutions  
Cloud infrastructure



Premier

**Business Partner**  
Solution provider

Container platform specialist



Advanced

**Training Partner**

Authorized reseller



What connects us all?  
Open-Source at the core!

Angelo Jacobs -  
Business Development  
Manager

# Piros Partnerships

Piros.

# Piros in numbers

Piros.

18 YEARS RED HAT PARTNERSHIP

7 YEARS ZABBIX PARTNERSHIP

+ 57% GROWTH IN 2024

+ 40 CERTIFIED RED HAT & ZABBIX ENGINEERS

+ 150 ENGINEERS IN THE CRONOS RED HAT COMMUNITY

+ 11.000 DE CRONOS GROEP COLLEAGUES

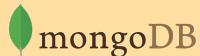
+ 600 COMPETENCE CENTERS



## Your Partner for Open Source Databases

Passionate  
Solution Driven  
Customer Focused

### Expertise & Partners



Official Reselling  
Partner

### Our Services:

- **Consultancy:** Expert **DBA support**: short-term help, a long-term partnership, or specialized training.
- **Project Work:** **Migrations**, Critical upgrades, Performance Tuning, Automation, **Security Assessment**, Brainstorm Sessions,...
- **Managed Services:** **8/18 or 24/7** - We provide proactive support, monitoring, and incident response.



# Red Hat Openshift and EDB Partnership



# Red Hat OpenShift with EDB

Tom D'Hont  
Senior Account Solution  
Architect - Ecosystem

# Red Hat is a Leader in the 2024 Gartner® Magic Quadrant™: Container Management

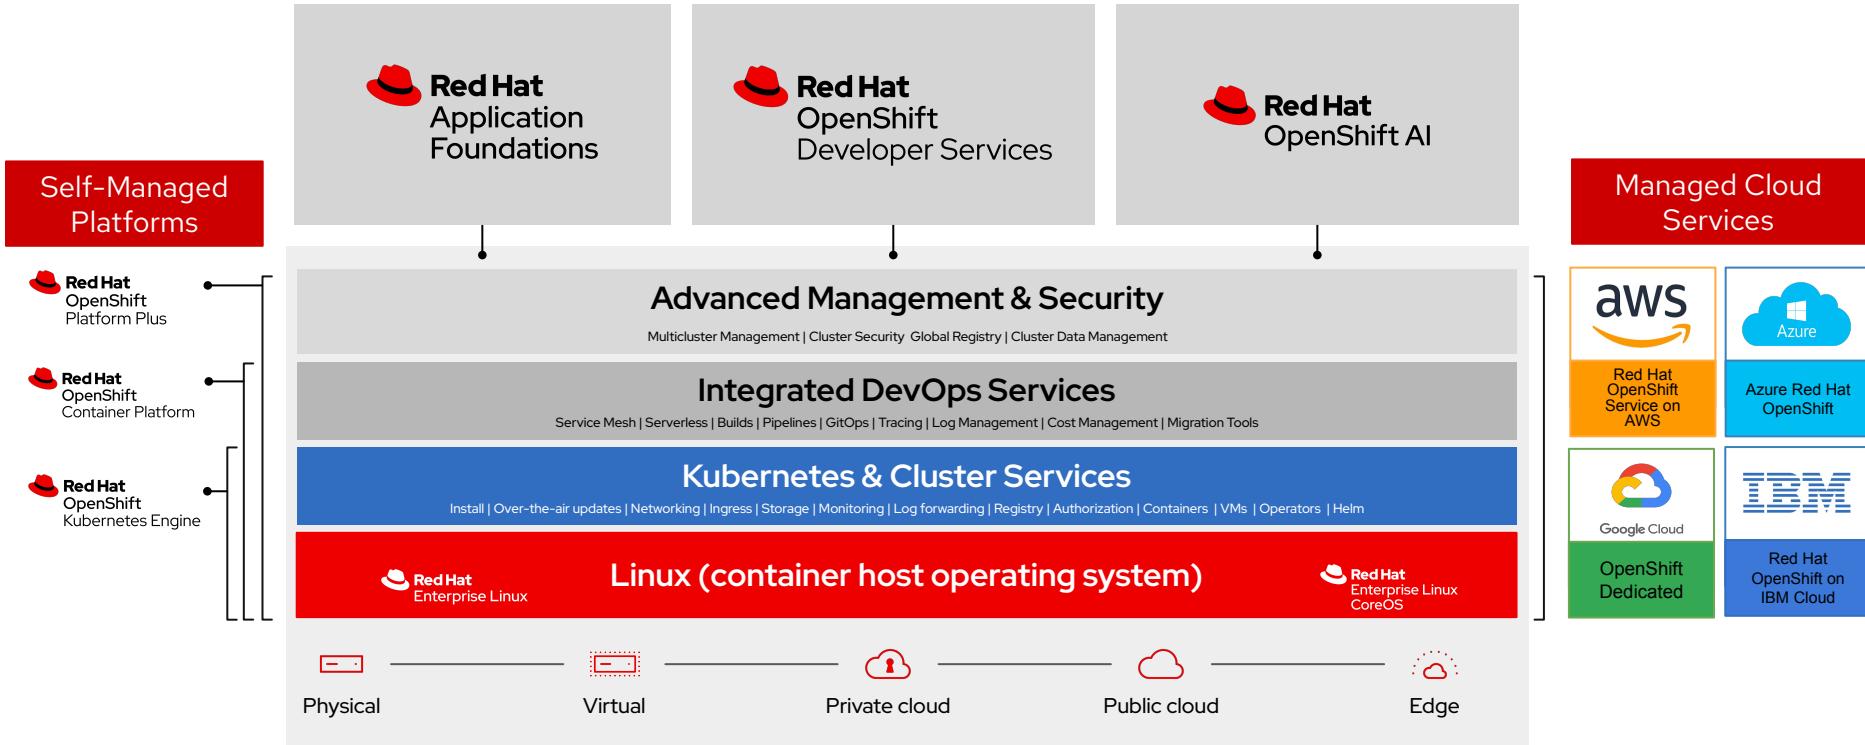
Figure 1: Magic Quadrant for Container Management



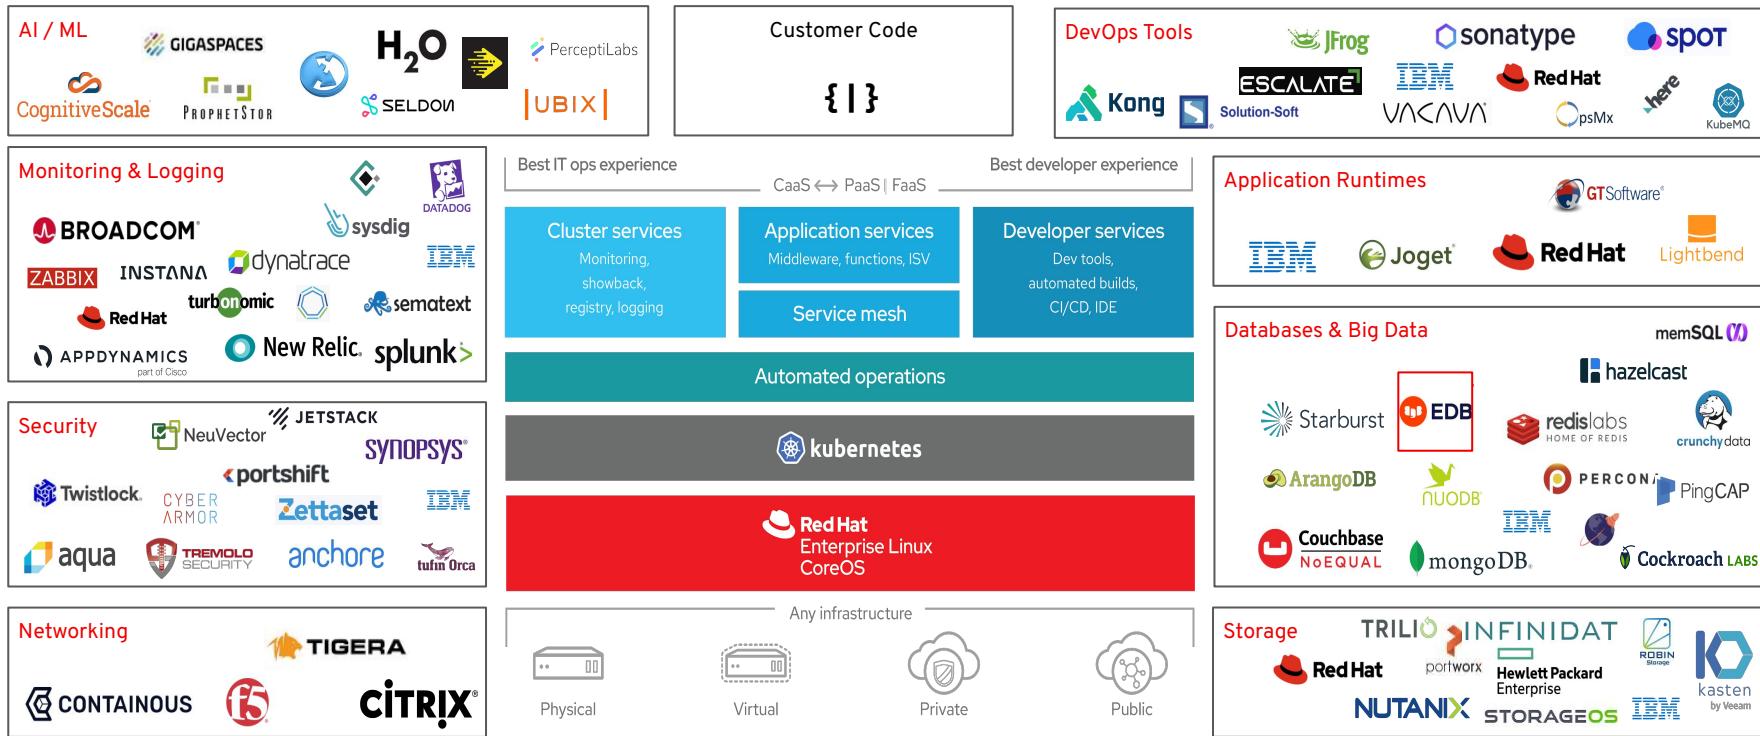
Gartner

Source: Gartner, "Magic Quadrant for Container Management," September 2024.

# Hybrid Cloud Application Platform



# Red Hat open hybrid cloud platform with ISV ecosystem



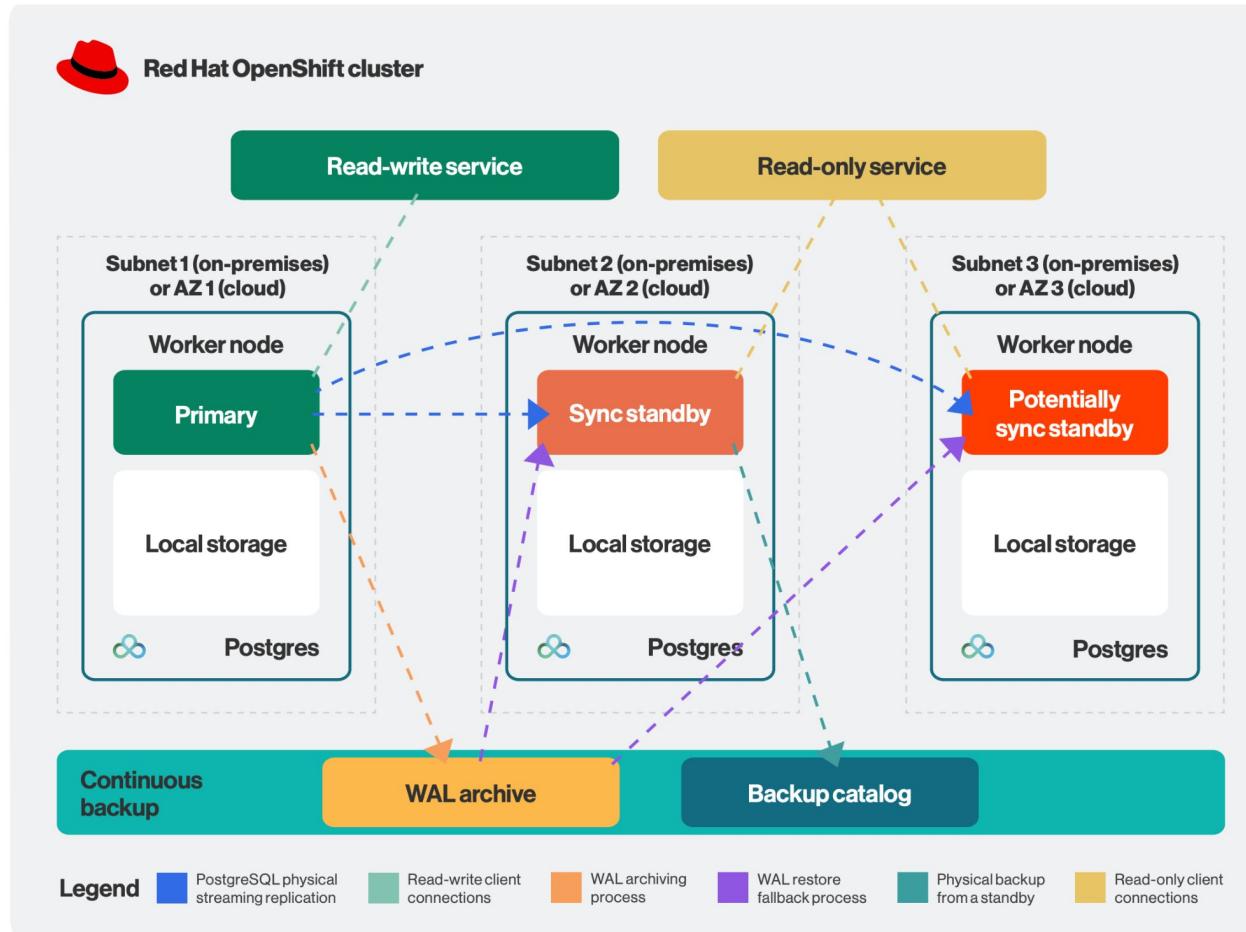
# Why Red Hat OpenShift for EDB: operator certification

The screenshot shows the Red Hat Ecosystem Catalog interface. At the top, there's a navigation bar with the Red Hat logo, 'Red Hat Ecosystem Catalog', and links for 'Solutions', 'Products', 'Artifacts', and 'Partners'. A search bar is also present. Below the navigation, the URL 'Home > Software > All software results > Containerized applications' is visible. The main content area features the 'EDB Postgres for Kubernetes' operator, which is certified. The page includes tabs for 'Overview', 'Resources', 'Certifications' (which is selected and highlighted with a red border), 'Deploy & use', and 'FAQs'. Below the tabs, there's a section titled 'Certifications' with a link to 'Learn about Red Hat Certification and Partner Validation'. Under 'Certified components', there's a table listing the operator with details like 'Image type' (Container image), 'Version' (1.25.1), and 'Architecture' (amd64). The table also shows the publisher as 'EnterpriseDB' and the status as 'Published last month'.

EDB Postgres for Kubernetes is a certified Level 5 Operator for Red Hat OpenShift

- ▶ This is designed to streamline Day 2 operations of PostgreSQL databases
- ▶ Enhanced Database Management
- ▶ Supports point-in-time recovery (PITR)
- ▶ Ensures robust data protection and recovery options
- ▶ Integration with business continuity solutions such as Red Hat OpenShift API for Data Protection (OADP) and Veeam Kasten, Trilio, Portworx Backup, IBM Fusion, and others

# Why Red Hat OpenShift for EDB : reference architecture



# EDB on OpenShift use cases

- ▶ Cloud-Native Database Deployment
- ▶ Database as a Service (DBaaS)
- ▶ High Availability and Disaster Recovery (HA & DR)
- ▶ DevOps and Continuous Integration/Continuous Deployment (CI/CD)
- ▶ Microservices and Application Modernization
- ▶ Move from VMWare to OpenShift
- ▶ Data Security and Compliance (using **TDE** and Advanced Security provided by EPAS)
- ▶ Hybrid and Multi-Cloud Deployments
- ▶ Multi-Tenant Applications (isolation)



# Euro Information

## Company profile

Euro-Information is the fintech company of the Crédit Mutuel group. Euro-Information manages the IT systems of 16 federations of Crédit Mutuel as well as those of CIC and of all the financial, insurance, property, consumer credit, private banking, financing, telephony and technological subsidiaries.

## Problem

- Fast database deployment
- Adopt a supported and secure Open Source platform
- Onprem DBaaS
- Align to in-house RDBMS standardization

## Solution

- Use Postgres capabilities to build and maintain local applications
- Use Red Hat OpenShift platform to accelerate the provisioning of databases and applications

## Results

- Applications running with PostgreSQL databases in a centralized environment
- Massive reduction of TCO of database service operations



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL
- EPAS

- EDB considerably reduces IT costs associated with database maintenance.
- 280 cores: Enterprise Plan + Production Support

## Summary

Use Case

On prem DBaaS ([in Production](#))

Workload

Transactional

Application Name

All internal Postgres applications

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes



# La Poste

## Company profile

La Poste is a postal service company in France, operating in Metropolitan France, the five French overseas departments and regions and the overseas collectivity of Saint Pierre and Miquelon. Under bilateral agreements, La Poste also has responsibility for mail services in Monaco through La Poste Monaco and in Andorra alongside the Spanish company Correos.

## Problem

- Provide a database HA solution for Ansible Automation Platform (AAP)
- Database must be in HA and DR

## Solution

- Use EDB Postgres for Kubernetes to provide a HA and DR solution for PostgreSQL databases
- Deploy in 2 OpenShift clusters our operator

## Results

- La Poste developer can use their internal 'La Post Service Portal' to provision more than 64 backends.
- Reduce risk deploying EDB solutions.



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL



- EDB considerably reduces IT costs associated with database maintenance.
- 12 Cores: Standard Plan + Premium Support

## Summary

Use Case

On prem DBaaS with HA and DR  
[\(In Production\)](#)

Workload

Transactional

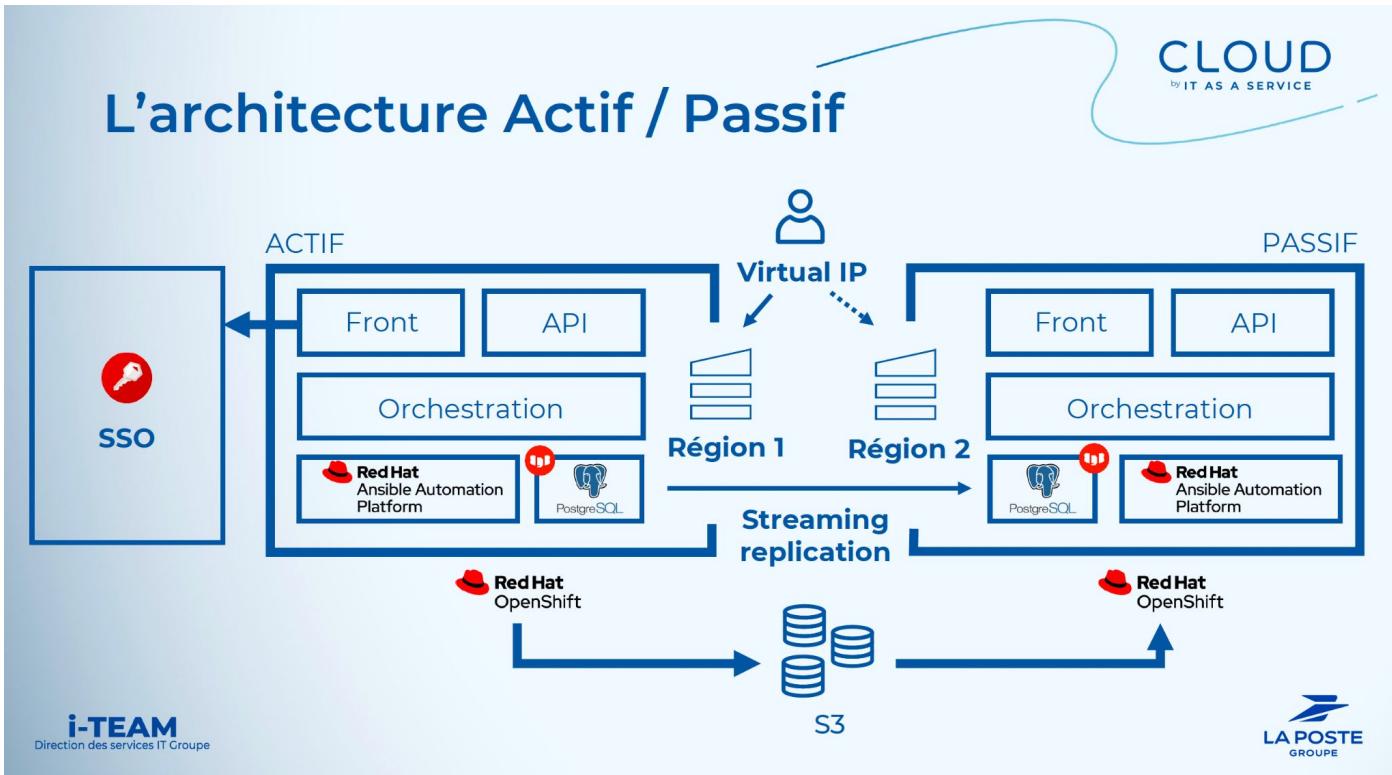
Application Name

Portail XaaS

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes

# La Poste Architecture



# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions.

Award-winning support, training, and consulting services make

Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)

# Introduction to CloudNativePG and EDB





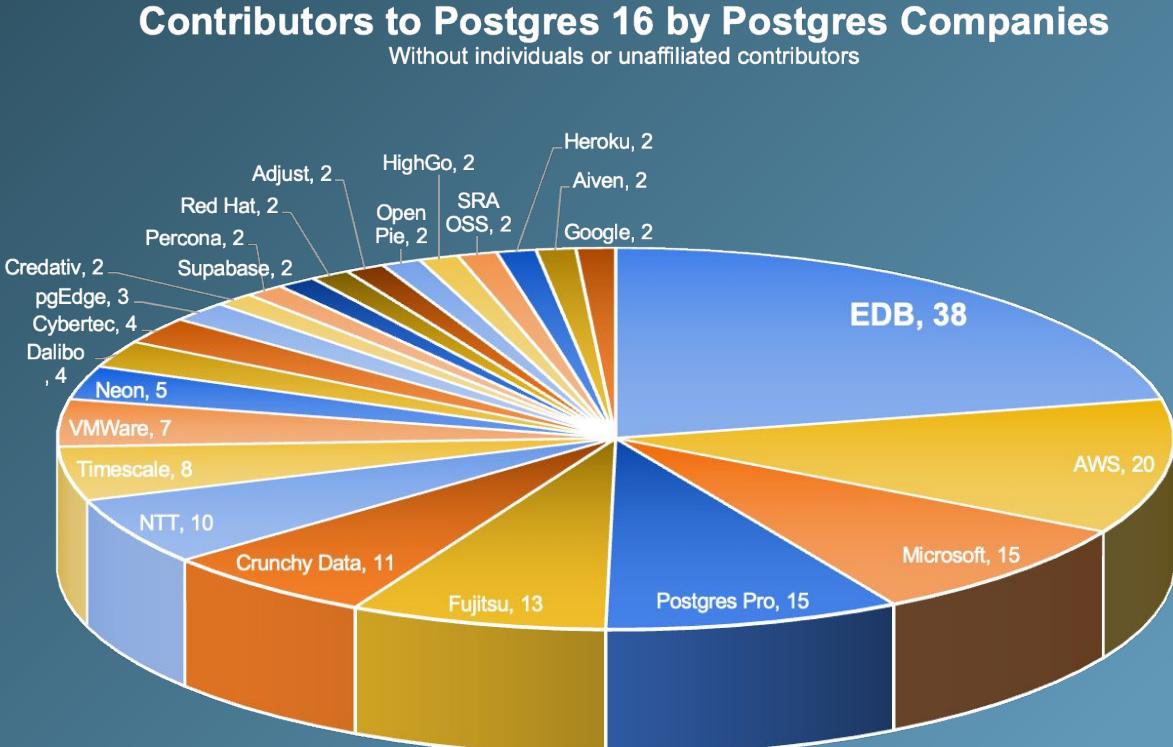
# 20+ years of Postgres innovation & adoption

- Number one contributor to Postgres, fastest-growing and most loved Database in the world
  - 2 Core Team members, 7 Committers, 9 Major Contributors, 10 Contributors, #1 site for desktop downloads
- Over 700 employees in more than 30 countries
- EDB Postgres AI
  - The industry's first platform that can be deployed as cloud, software or physical appliance
  - Secure, compliant and enterprise grade performance guaranteed



# Large Developer Community

- 407 + contributors to Postgres 16
- 111 companies actively contributing to Postgres 16



# Postgres has won the database race



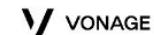
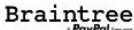
Stack Overflow Survey 2023/2024



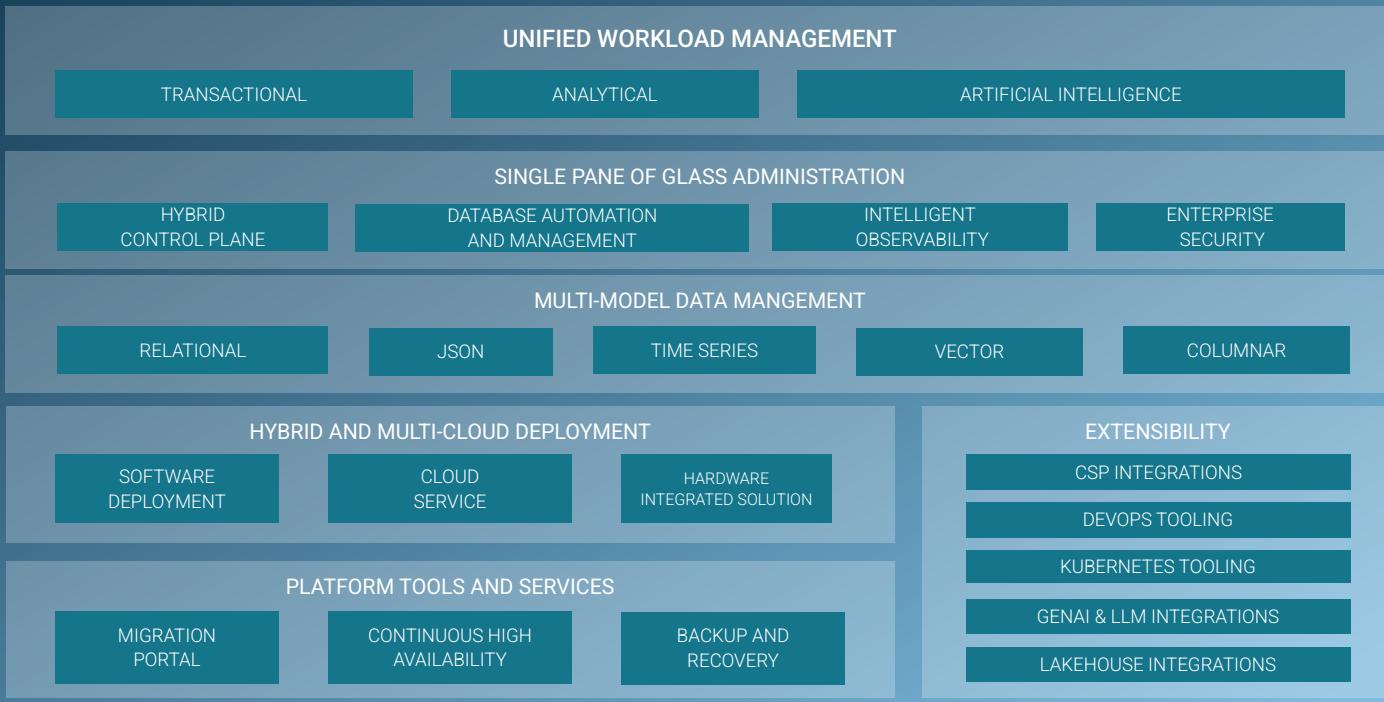
BANKING FINANCIAL

TECHNOLOGY

TELCO



# EDB POSTGRES AI PLATFORM





FERRARI SF-24



Chance to WIN  
LEGO Technic 42207  
Ferrari!

# CNPG Operator: Reference Architecture and functionalities

# Kubernetes timeline

- 2014, June: Google open sources Kubernetes
- 2015, July: Version 1.0 is released
- 2015, July: Google and Linux Foundation start the CNCF
- 2016, November: The operator pattern is introduced in a blog post
- 2018, August: The Community takes the lead
- 2019, April: Version 1.14 introduces **Local Persistent Volumes**
- 2019, August: EDB team starts the Kubernetes initiative
- 2020, June: we publish this blog about benchmarking local PVs on bare metal
- 2020, June: Data on Kubernetes Community founded
- 2021, February: EDB Cloud Native Postgres (CNP) 1.0 released
- 2022, May: **EDB donates CNP** and open sources it under CloudNativePG
- 2025, January: CloudNativePG was recognized as an official **#CNCF** project



# Enabling the same PostgreSQL everywhere

From self-managed to fully managed DBaaS in the Cloud

- Same applications
- Faster innovation
- Performance and scalability
- Stability, security and control
- Seamless integration
- Obsolescence



Private

Hybrid

Multi-cloud

Public



Bare Metal



Virtual Machines



Containers



# A kubernetes operator for Postgres



Kubernetes adoption is rising and it is already the de facto **standard** orchestration tool



PostgreSQL clusters “management the kubernetes way” enables many cloud native usage patterns, e.g. spinning up, disposable clusters during tests, one cluster per microservice and one database per cluster

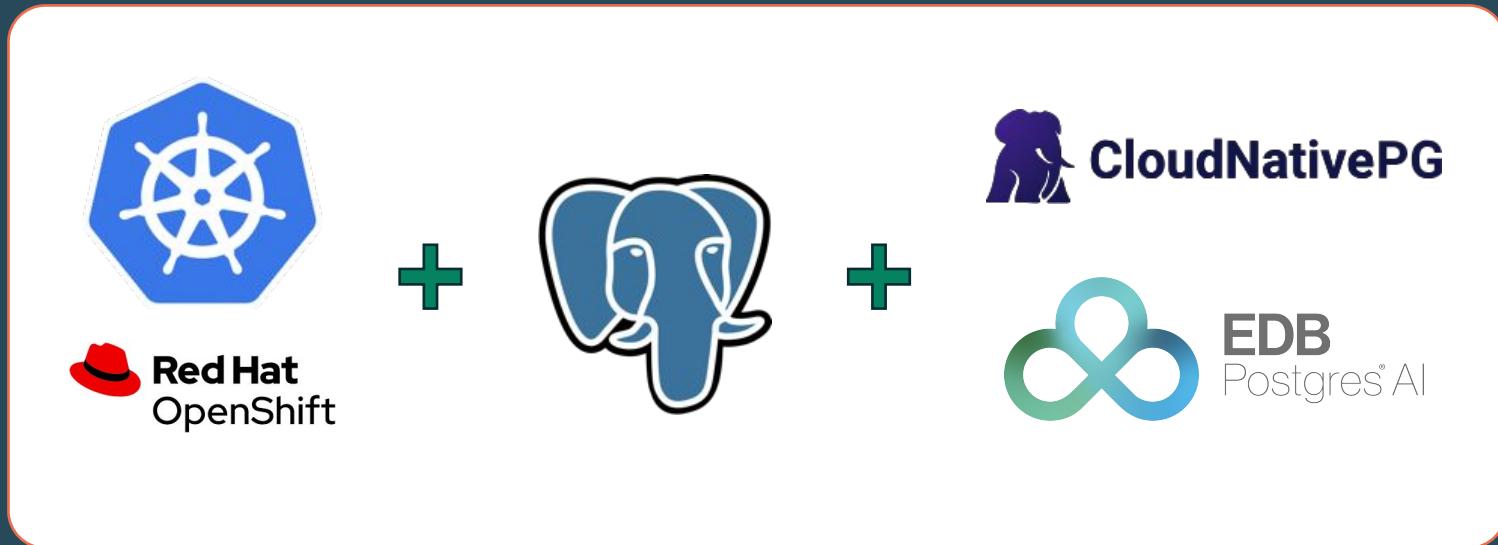


CNPG tries to encode years of experience managing PostgreSQL clusters into **an Operator which should automate all the known tasks** a user could be willing to do

Our PostgreSQL operator must simulate the work of a DBA



# Win Technology



## Autopilot

It automates the steps that a human operator would do to deploy and to manage a Postgres database inside Kubernetes, including automated failover.



## Security

A grayscale photograph of a security guard from behind. The guard is wearing a light-colored jacket with the word "SECURITY" printed in large, bold, white capital letters on the back. He is holding a dark walkie-talkie up to his ear with his right hand. The background is a plain, light-colored wall.

CloudNativePG is secured by default.





It doesn't rely on statefulsets and uses its own way to manage persistent volume claims where the PGDATA is stored.

## Data persistence



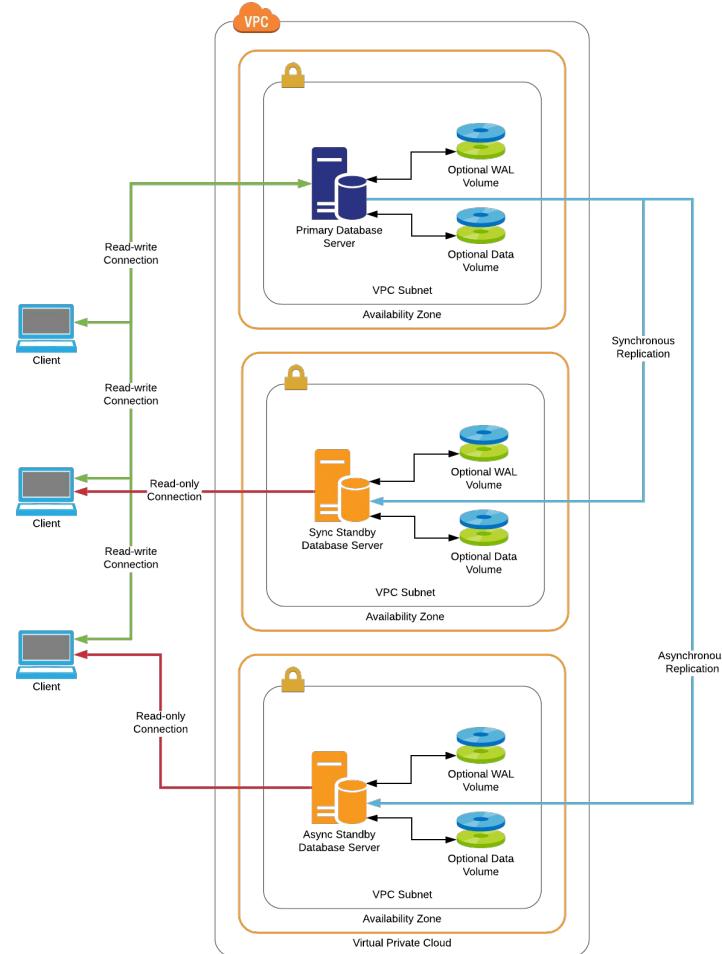
## Designed for Kubernetes

It's entirely declarative, and directly integrates with the Kubernetes API server to update the state of the cluster — for this reason, it does not require an external failover management tool.



# Imperative vs Declarative

- Create and configure VMs
- Create a PostgreSQL 13 instance
- Configure for replication
- Clone a second one
- Set it as a replica
- Clone a third one
- Set it as a replica
- Configure networking
- Configure security
- etc.



# Convention over configuration

Declarative - simple to install, simple to maintain

There's a PostgreSQL 17 cluster with 2 replicas:

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: myapp-db
spec:
  instances: 3
  imageName: quay.io/enterprisedb/postgresql:17

storage:
  size: 10Gi
```



# Features

Deployment	Administration	Backup & Recovery	Monitoring	Security	High Availability
Kubernetes operator	Single node	Backup	Prometheus	TDE	Switchover
Kubernetes plugin	Cluster (Multi node)	Recovery	Grafana dashboards	Certificates	Failover
EDB Postgres (EPAS)	PostgreSQL configuration	PITR	Postgres Enterprise Manager	Data redaction	Scale out / scale down
PostGIS	Pooling	Volume Snapshots	Logging	Password management	Minor / Major updates



# Decision-making for choosing the deployment platform



# When to choose Kubernetes over VMs?

**01** | Cloud Native Applications that already run in Kubernetes

**02** | Scalable, replicated databases

**03** | Applications requiring automated failover and self-healing

**04** | Teams skilled in Kubernetes who want a unified infrastructure



# Advantage of deploying Postgres Databases in Kubernetes

## Automation & Orchestration

- 01 |
- Self-healing
  - Automated scaling
  - Rolling updates

## Self-healing

- 02 |
- Best resource utilization
  - Dynamic Resource allocation

## Rolling updates

- 03 |
- Cloud-agnostic
  - Consistent deployment

## Service discovery & networking

- 04 |
- Built-it service discovery
  - Load Balancing

## Automated backups and disaster recovery

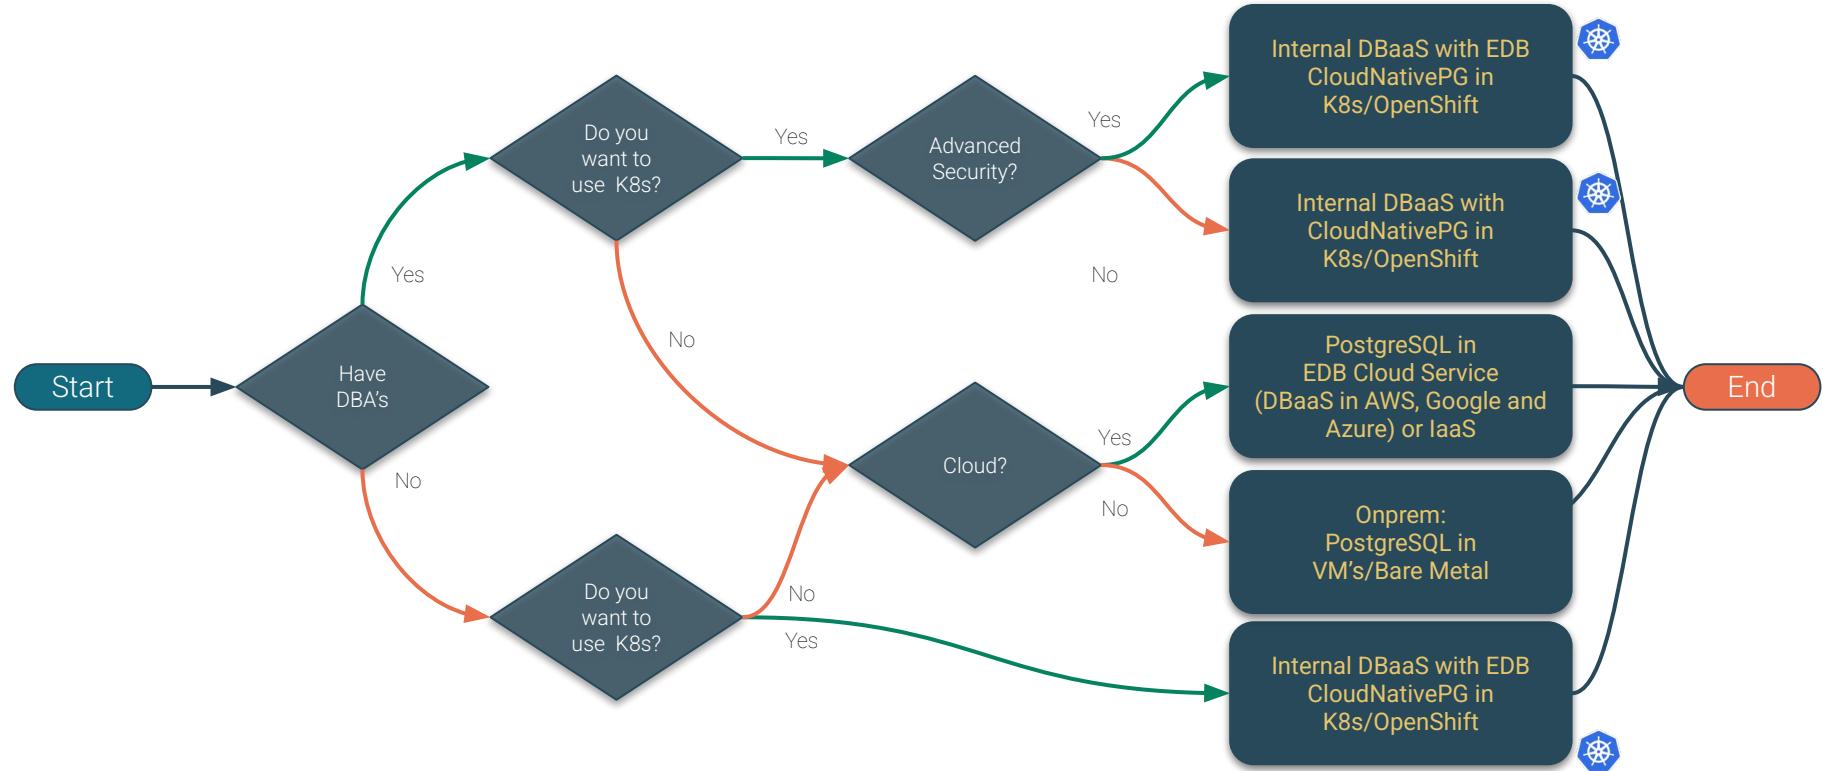
- 05 |
- Automated backups
  - Multi-region failover

## Security & access control

- 06 |
- RBAC
  - Secret management



# Decision-making for choosing the deployment platform

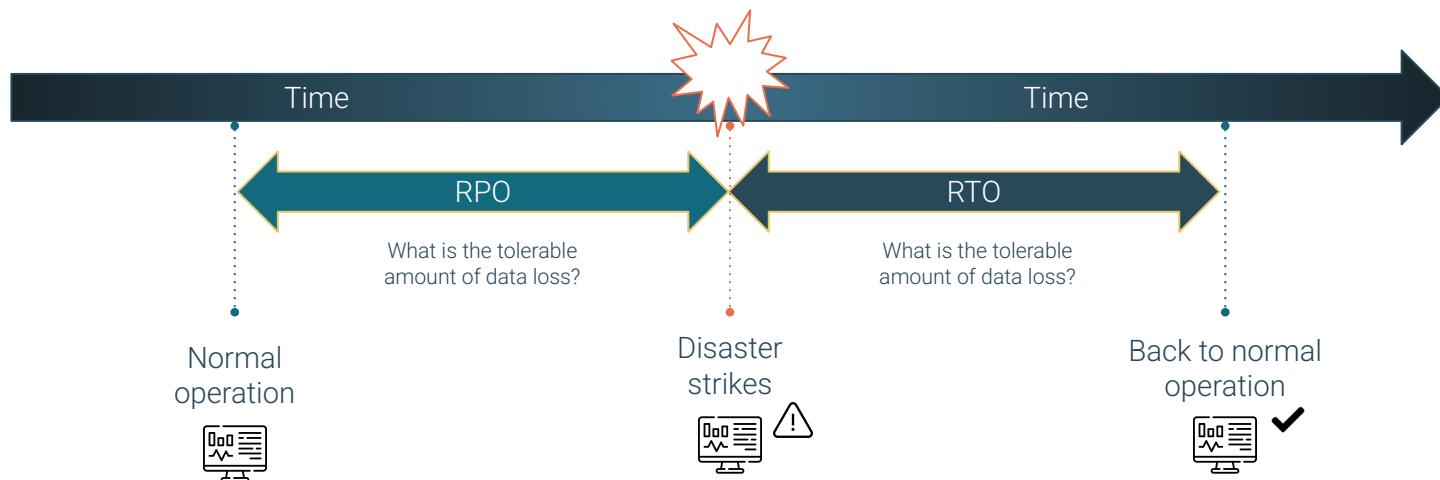


# Architectures



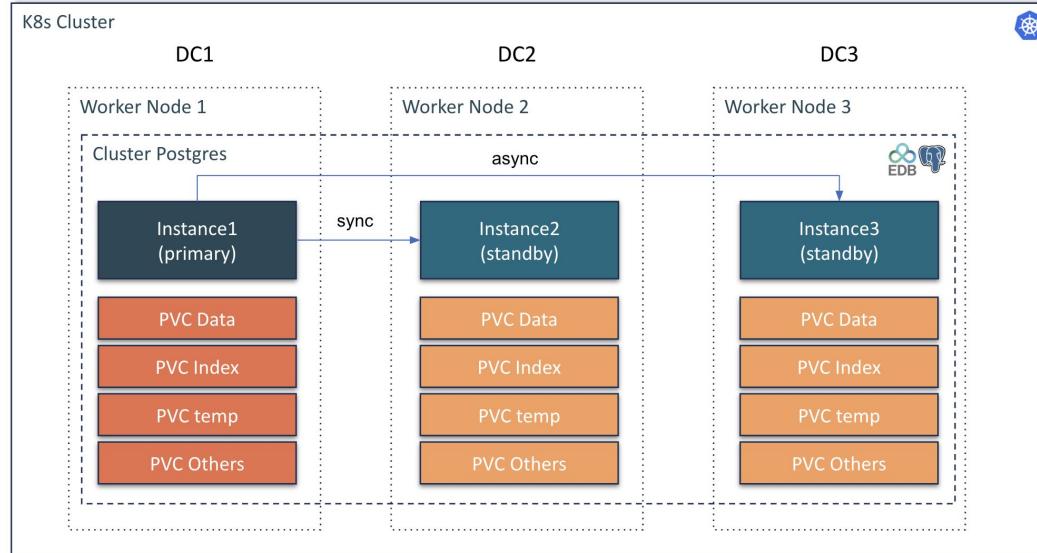
# Concepts

- Recovery Point Objective (**RPO**) and Recovery Time Objective (**RTO**) are key concepts in disaster recovery and business continuity planning, particularly related to data loss and system downtime.



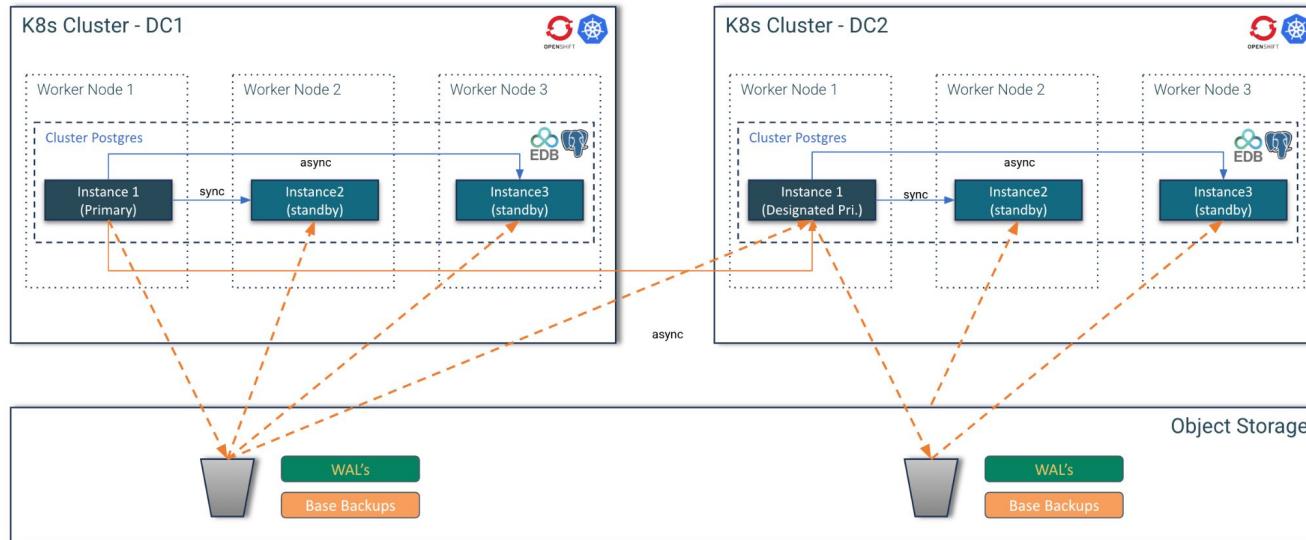
# Red Hat Recommendation

Red Hat recommend stretched clusters ONLY when latencies don't exceed 5 milliseconds (ms) round-trip time (RTT) between the nodes in different locations, with a maximum RTT of 10 ms.

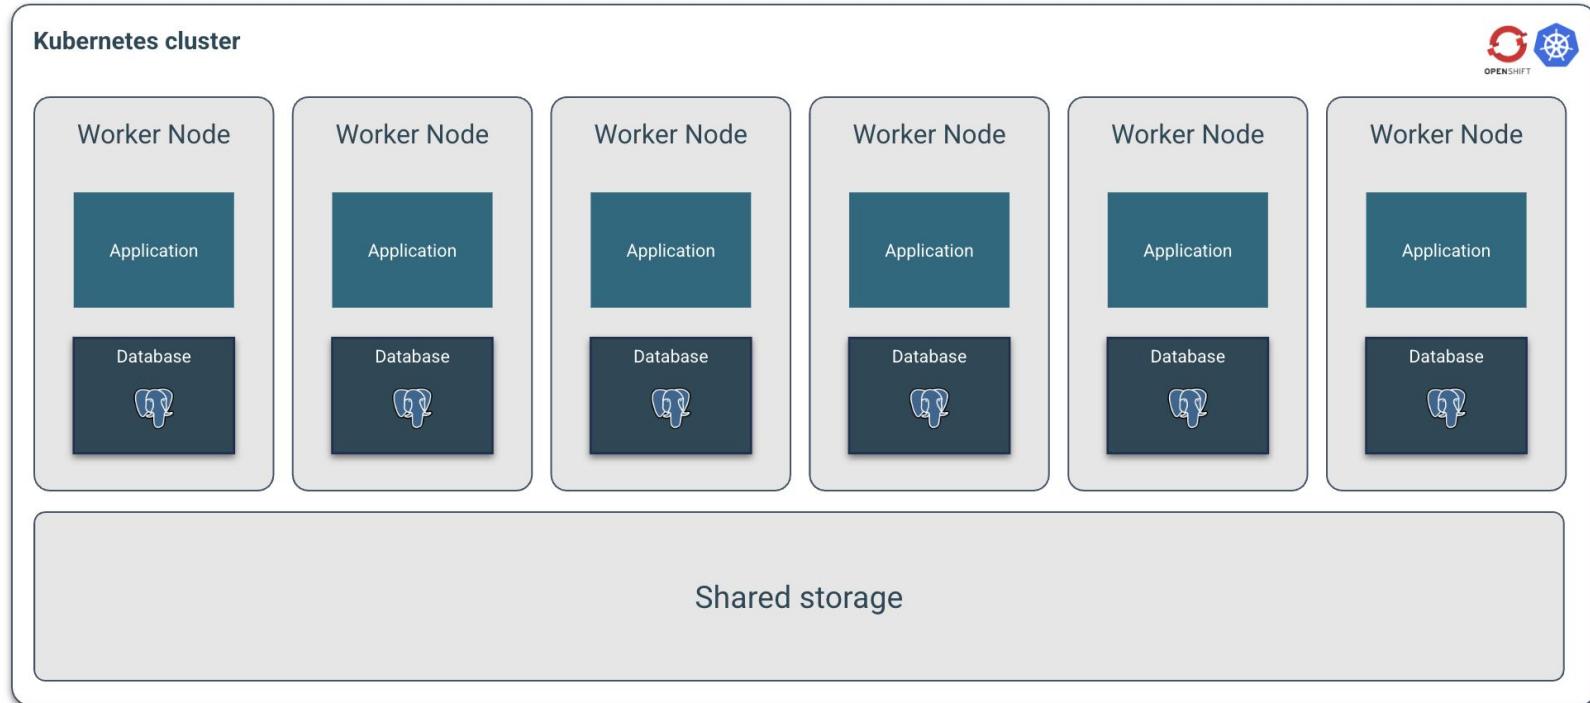


# Two separate single data center Kubernetes clusters

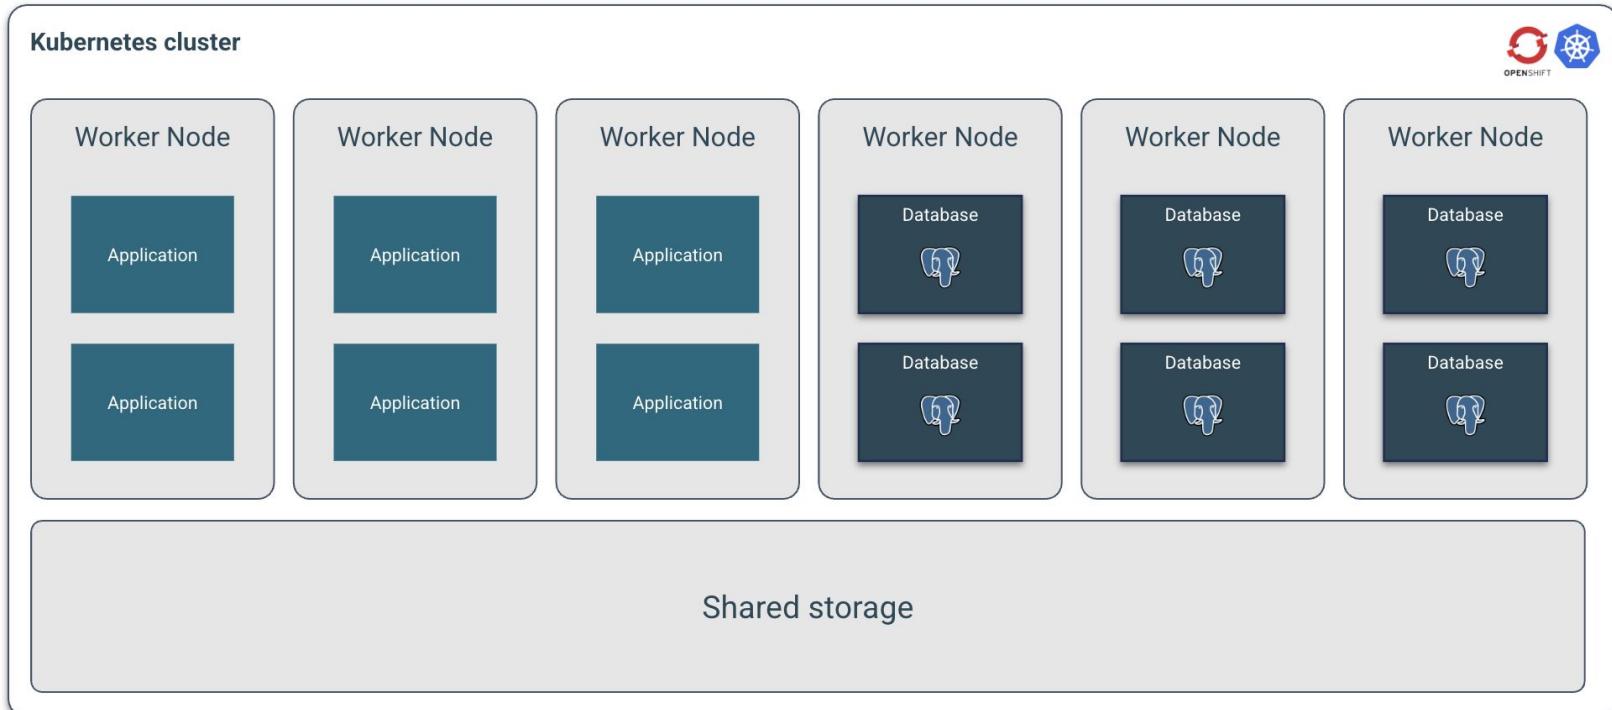
In case you cannot go beyond two data centers and you end up with two separate Kubernetes clusters, don't despair.



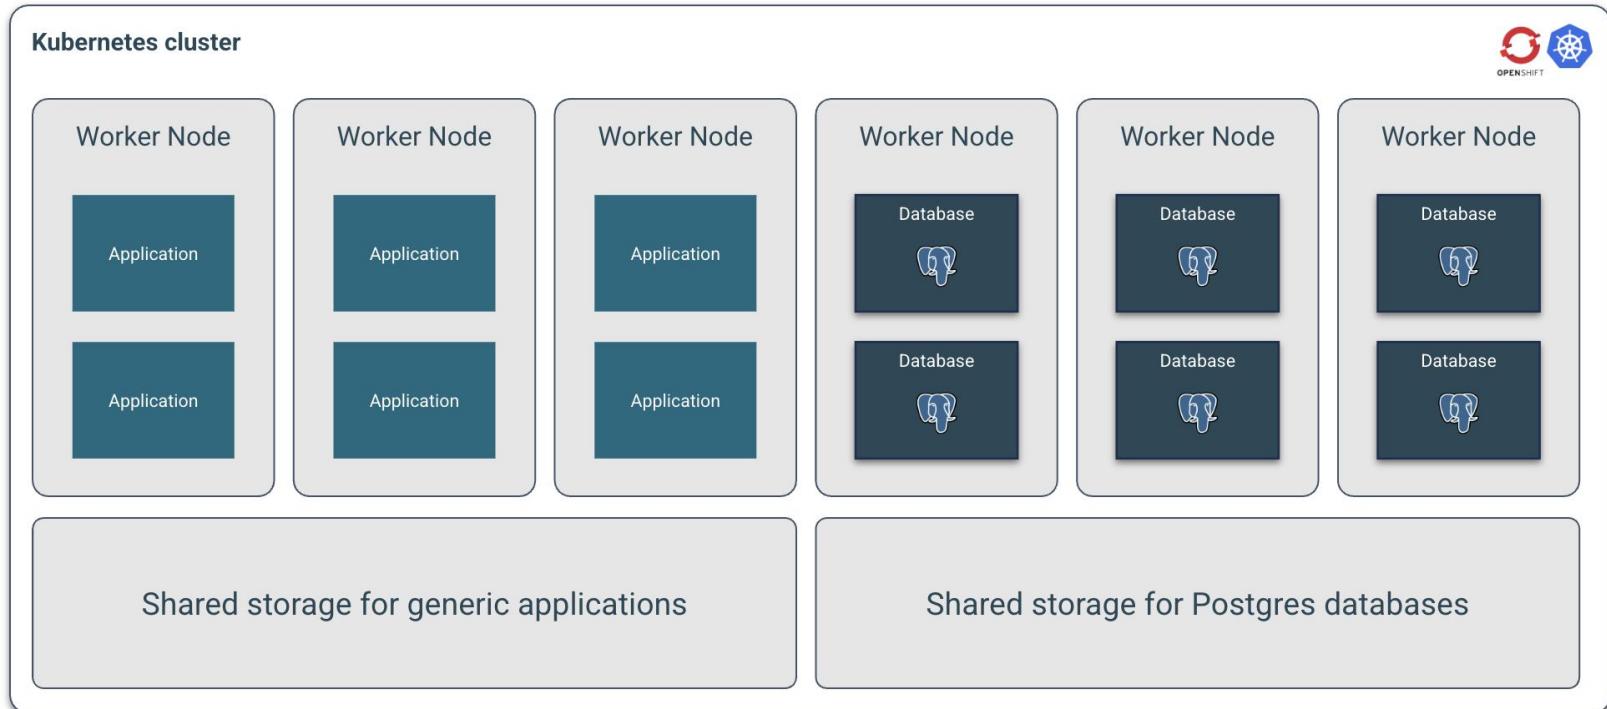
# Shared workload, shared storage



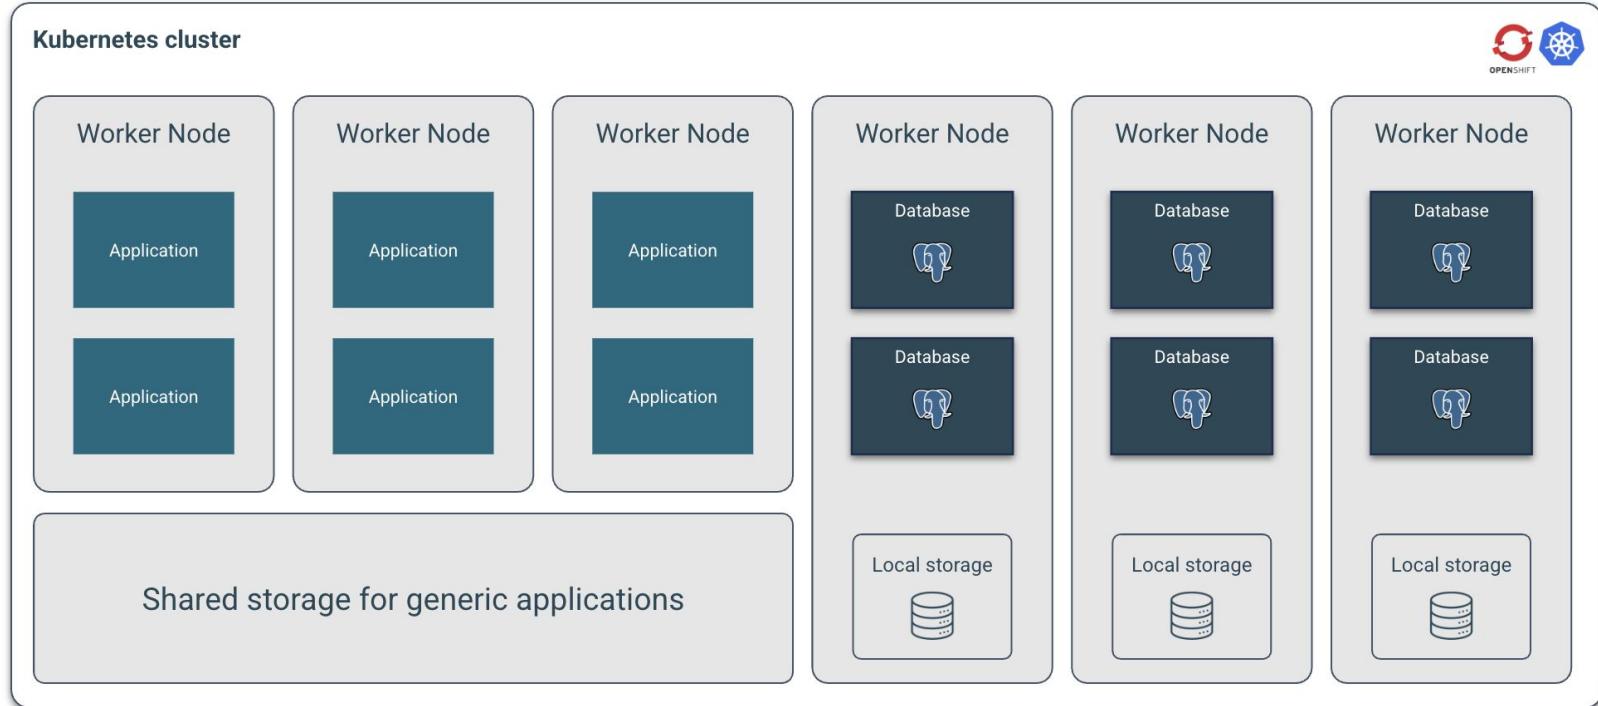
# Shared workload, shared storage



# Shared workload, shared storage



# Shared workloads, local storage



# Recommended architectures

<https://www.cncf.io/blog/2023/09/29/recommended-architectures-for-postgresql-in-kubernetes/>



## Recommended architectures for PostgreSQL in Kubernetes

By Gabriele Bartolini

September 29, 2023

Member post by Gabriele Bartolini, VP of Cloud Native at EDB

"You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM", [tweeted Kelsey Hightower just a few months ago](#). Quite the opposite from what the former Google engineer and advocate said back in 2018 on Twitter: "Kubernetes supports stateful workloads; I don't."



Kelsey Hightower @kelseyhightower

You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM. The biggest challenge is understanding that rubbing Kubernetes on Postgres won't turn it into Cloud SQL.

Truth is that I agree with him now as much as I agreed with him back then. At that time, the holistic offering of storage capabilities in Kubernetes was still immature (local persistent volumes would become GA only the year after), the operator pattern – which in the meantime has proven to be crucial for stateful applications like databases – was yet to become widely accepted, and the [Data on Kubernetes Community](#) was more than two years away (second half of 2020).

Nowadays, the situation is completely different. And I am sure that many people who've worked hard in the last few years to bring stateful workloads in Kubernetes agree with me that Kelsey's recent powerful words will contribute to reversing the public perception and facilitate our mission - provided we keep doing great.



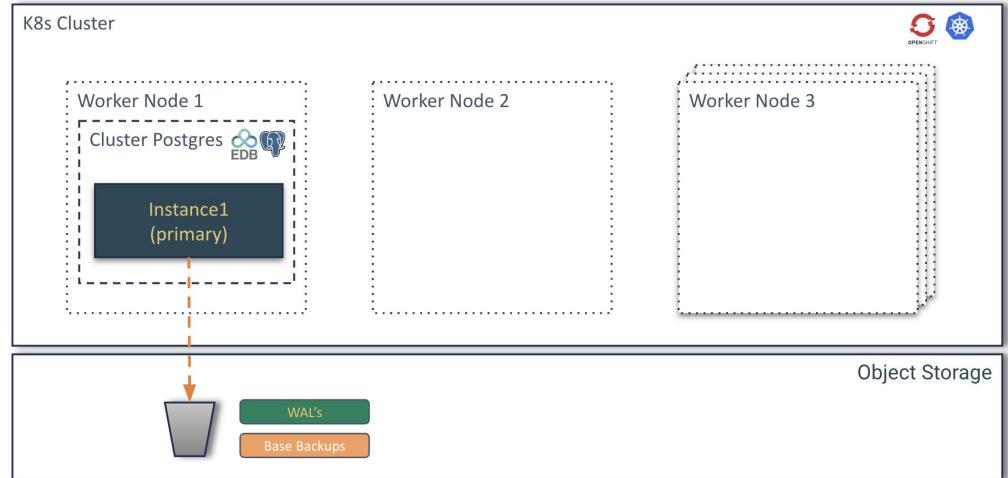
# Use Cases



# Use case 1 architecture

A single database is the simplest setup, involving one instance of a database server.

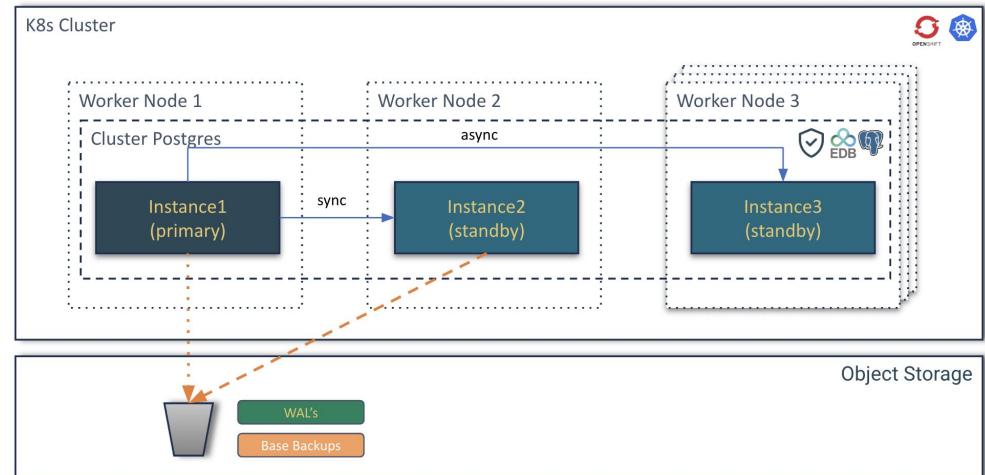
- Development and testing environments
- Small applications with low traffic
- Non-critical data analysis
- Applications with high tolerance for downtime
- Cost-sensitive projects



# Use case 2 architecture

An HA database setup aims to minimize downtime by having redundant components. If one component fails, another takes over automatically or with minimal intervention. This usually involves techniques like clustering, replication, or mirroring within the same data center or availability zone.

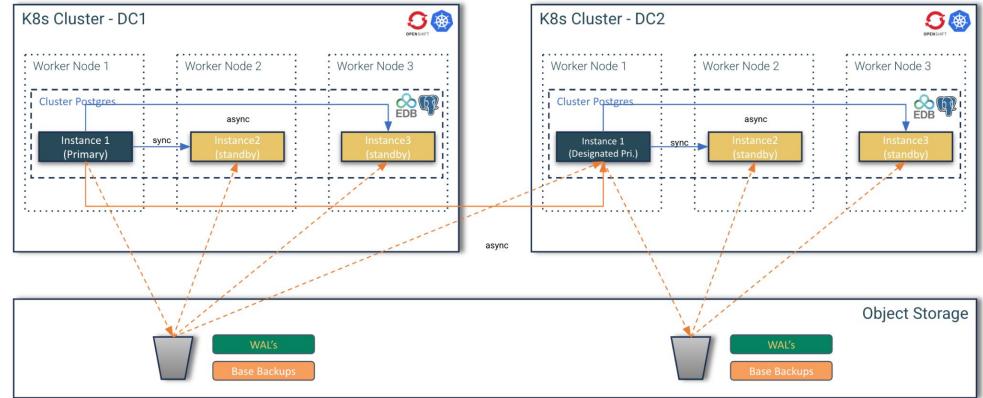
- Business critical Applications
- Applications with stringent SLAs
- Real-time systems
- Improving user experience
- Minimizing planned downtime



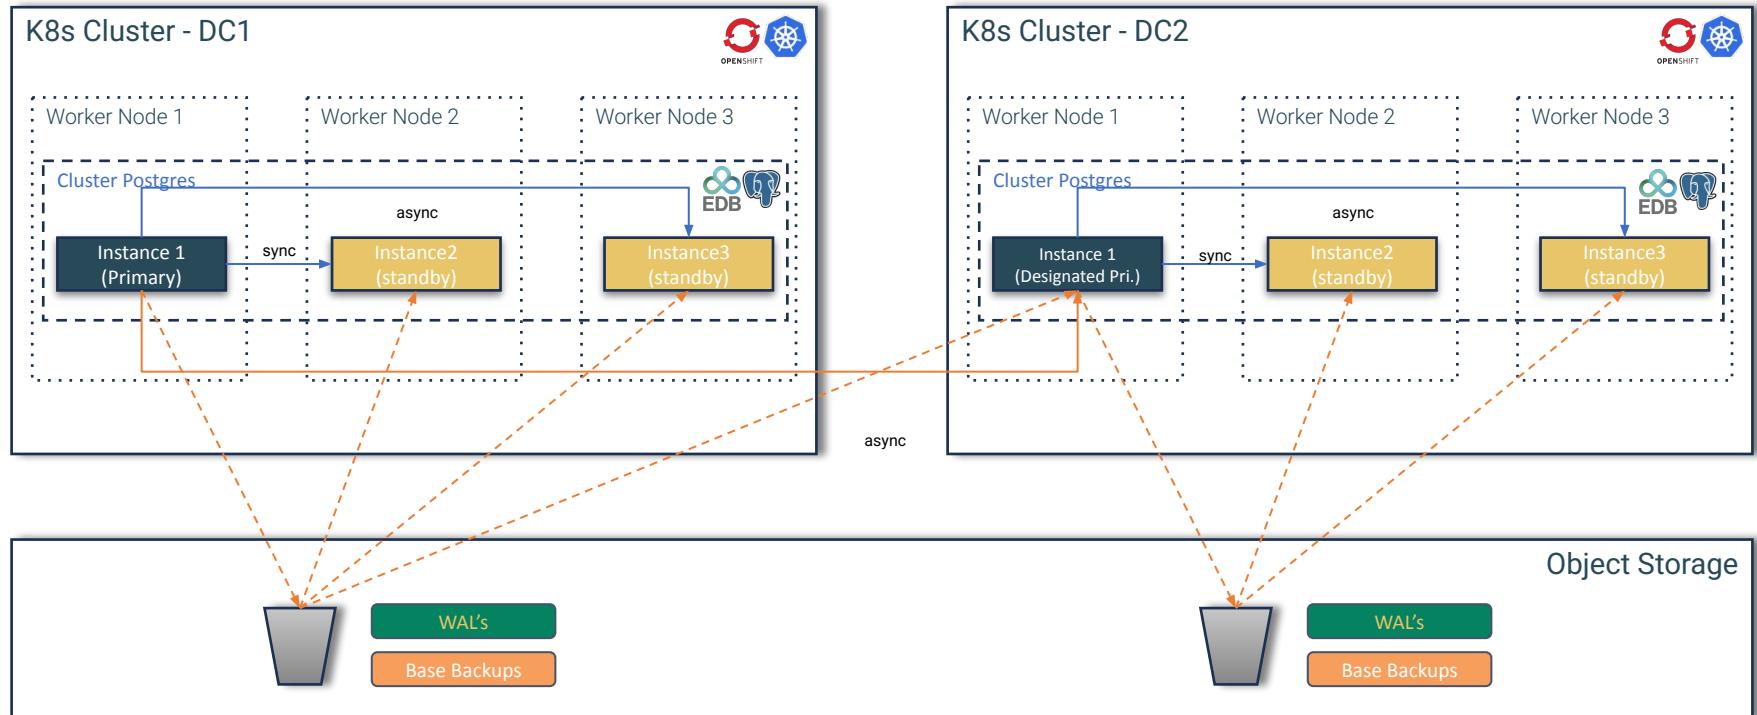
# Use case 3 architecture

A DR database setup focuses on protecting data and ensuring business continuity in the event of a large-scale disaster affecting an entire data center or region (e.g., natural disasters, power outages, cyberattacks). This typically involves replicating data to a geographically separate location.

- Regulatory compliance
- Protecting against catastrophic data loss
- Ensuring business continuity for mission-critical systems



# Use case 3 architecture



# Interactive session

# It's time to go hands-on!



Hand-on documentation



Download this presentation

<https://tinyurl.com/3j7cbjh3>



# Links:

## Openshift Console:

<https://console-openshift-console.apps.cluster-m6pll.m6pll.sandbox3121.opentlc.com>

## Users:

name: user2..user40  
Password: edb-workshop

## Devspaces url:

<https://devspaces.apps.cluster-m6pll.m6pll.sandbox3121.opentlc.com/>

## Short url to Devspaces:

<https://tinyurl.com/yy9dswmk>

## Minio:

UI: <https://minio-ui-default.apps.cluster-m6pll.m6pll.sandbox3121.opentlc.com>  
API: <https://minio-api-default.apps.cluster-m6pll.m6pll.sandbox3121.opentlc.com>

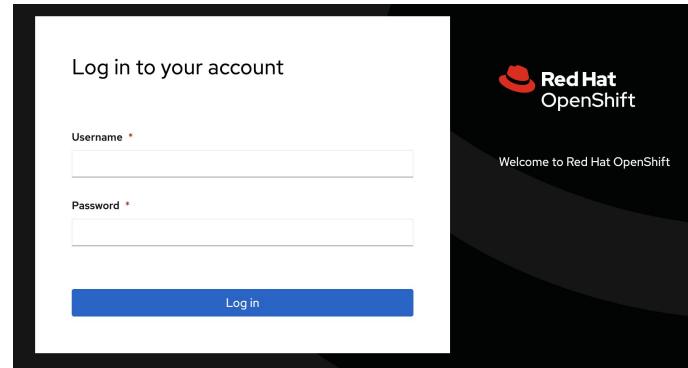
User: minio  
Password: edb-workshop



Open the following URL in your browser:

 <https://red.ht/edb-cph25>

<https://tinyurl.com/yy9dswmk>



Username and password provided to you

## Authorize Access

openshift-operators-client is requesting permission to access your account (user)

### Requested permissions

user:full

Full read/write access with all of your permissions

Includes any access you have to escalating resources like secrets

You will be redirected to <https://devspaces.apps.cluster-52d72.dynamic.redhatworkshops.io/oauth/callback>

[Allow selected permissions](#)

[Deny](#)

Press “Allow selected permissions”

## Select a Sample

Select a sample to create your first workspace.

workshop x 1 item



Postgres on OpenShift

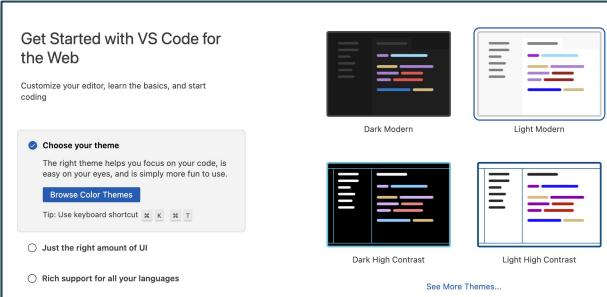
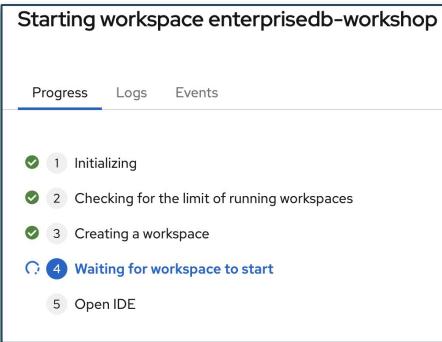
Workshop

CloudNativePG on OpenShift

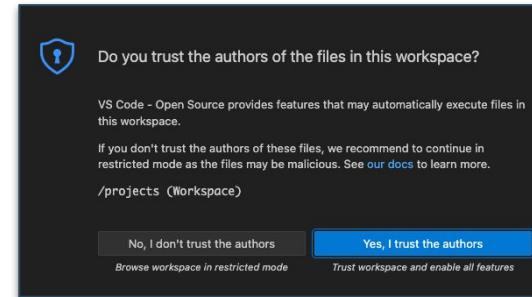
Workshop

In the Select a Sample section search for “Workshop” and click on the tile

## Environment information



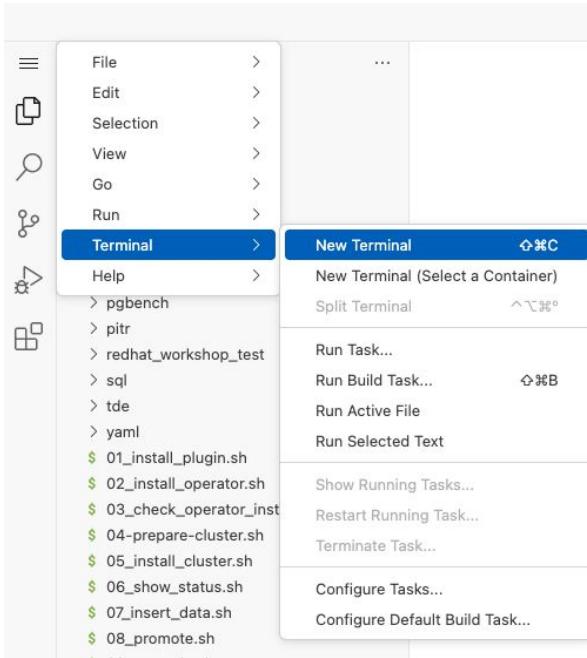
Select your theme



And trust the authors

```
$ config.sh ×  
workshop > $ config.sh  
1  #!/bin/bash  
2  
3  ./commands.sh  
4  
5  # Variables to be replaced  
6  export id=<user>                      # your name or id
```

Update your id with the user name you have been assigned



Open a new terminal

# Use case

# The environment



# Features shown during the demo

- Kubernetes plugin install
- Check the CloudNativePG operator status
- Postgres cluster install
- Insert data in the cluster
- Failover
- Backup
- Recovery
- Scale out/down
- Fencing
- Hibernation
- Monitoring
- Rolling updates (minor and major)

Deployment

High Availability

Administration

Monitoring

Backup and Recovery

Last CloudNativePG tested version is 1.25

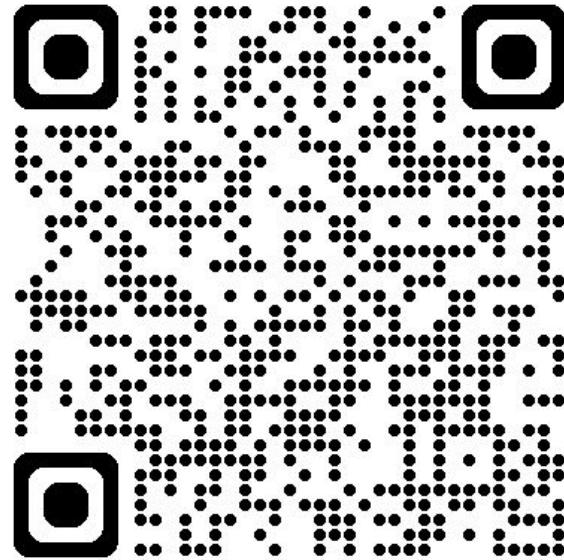


This demo is in



<https://github.com/sergioenterprisedb/edb-postgres-for-kubernetes-in-openshift>

<http://bit.ly/4duKxm7>



# Use case Plug-in installation



# The “cnp” plugin for kubectl

- The official CLI for CloudNativePG
  - Available also as RPM or Deb package
- Extends the ‘kubectl’ command:
  - Customize the installation of the operator
  - Status of a cluster
  - Perform a manual switchover (promote a standby) or a restart of a node
  - Issue TLS certificates for client authentication
  - Declare start and stop of a Kubernetes node maintenance
  - Destroy a cluster and all its PVC
  - Fence a cluster or a set of the instances
  - Hibernate a cluster
  - Generate jobs for benchmarking via pgbench and fio
  - Issue a new backup
  - Start pgadmin

**NOT NEEDED DURING WORKSHOP**  
**For illustrative purposes.**



```

Name: cluster-example
Namespace: default
System ID: 7100921006673293335
PostgreSQL Image: ghcr.io/cloudnative-pg/postgresql:14.3
Primary instance: cluster-example-2
Status: Cluster in healthy state
Instances: 3
Ready instances: 3
Current Write LSN: 0/C000060 (Timeline: 4 - WAL File: 000000040000000000000000C)

```

**NOT NEEDED DURING WORKSHOP**  
 For illustrative purposes.

#### Certificates Status

Certificate Name	Expiration Date	Days Left Until Expiration
cluster-example-replication	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-server	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-ca	2022-08-21 13:15:00 +0000 UTC	89.95

#### Continuous Backup status

```

First Point of Recoverability: 2022-05-23T13:37:08Z
Working WAL archiving: OK
WALs waiting to be archived: 0
Last Archived WAL: 00000004000000000000000B @ 2022-05-23T13:42:09.37537Z
Last Failed WAL: -

```

#### Streaming Replication status

Name	Sent LSN	Write LSN	Flush LSN	Replay LSN	Write Lag	Flush Lag	Replay Lag	State	Sync State	Sync Priority
cluster-example-3	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0
cluster-example-1	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0

#### Instances status

Name	Database Size	Current LSN	Replication role	Status	QoS	Manager Version
cluster-example-3	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0
cluster-example-2	33 MB	0/C000060	Primary	OK	BestEffort	1.15.0
cluster-example-1	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0



# Install CNPG plugin

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

- In the web terminal run the script 01\_install\_plugin.sh:

```
./01_install_plugin.sh
```

- Call the help for the CNPG Plugin, run:

```
kubectl-cnp help
```

## Try it for yourself

05:00



# Use case Operator installation



# Operator Installation demonstration

- Install the operator
- Check the installed CNP Operator in the console
- Discover the features of the Operator in the OpenShift environment
- Check the installed CNP Operator in the web terminal



NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

# Install the CNPG Operator and check it in the web terminal

- In the web terminal check the installation of the operator:

```
./02_install_operator.sh (will require admin privs on Openshift)
```

```
./03_check_operator_installed.sh
```

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.



# Check the installed CNPG Operator in the Openshift console

- In the OpenShift console navigate to:
  - -> Operators
  - -> Installed Operators
  - -> Click on the Operator installed in your namespace, for example: user1:

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

The screenshot shows the Red Hat OpenShift web interface. The left sidebar has a dark theme with the Red Hat logo and 'Red Hat OpenShift'. The 'Operators' section is expanded, and 'Installed Operators' is selected. The main content area shows a table titled 'Installed Operators'. The table has columns: Name, Managed Namespaces, Status, Last updated, and Provided APIs. One operator is listed: 'EDB Postgres for Kubernetes' (version 1.24.3 provided by EDB), which is managed across 'All Namespaces', has a status of 'Succeeded' (Up to date), was last updated on '21 May 2025, 15:21', and provides APIs for Backups, Cluster Image Catalog, Cluster, and Image Catalog. A note at the top of the table says: 'Installed Operators are represented by ClusterServiceVersions within this Namespace. For more information, see the [Understanding Operators documentation](#). Or create an Operator and ClusterServiceVersion using the [Operator SDK](#)'.

Name	Managed Namespaces	Status	Last updated	Provided APIs
EDB Postgres for Kubernetes 1.24.3 provided by EDB	All Namespaces	<span style="color: green;">● Succeeded</span> Up to date	21 May 2025, 15:21	Backups Cluster Image Catalog Cluster Image Catalog View 2 more...



# Discover the features of the Operator in the OpenShift environment

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

The screenshot shows the Red Hat OpenShift web interface. The left sidebar is titled "Administrator" and includes sections for Home, Operators (selected), OperatorHub, Installed Operators, Workloads, Networking, Storage, Builds, Compute, User Management, and Administration. The main content area is titled "Project: openshift-operators". It shows the "Installed Operators" section with "EDB Postgres for Kubernetes" selected. The "Operator details" page displays the following information:

- Provider:** EDB
- Created at:** 21 May 2025, 15:21
- Links:** EDB Postgres for Kubernetes (<https://www.enterprisedb.com/products/postgresql-on-kubernetes-ha-clusters-k8s-containers-scalable>)
- Documentation:** ([https://www.enterprisedb.com/docs/postgres\\_for\\_kubernetes/latest/](https://www.enterprisedb.com/docs/postgres_for_kubernetes/latest/))
- Maintainers:** Jonathan Gonzalez V. ([jonathan.gonzalez@enterprisedb.com](mailto:jonathan.gonzalez@enterprisedb.com)), Jonathan Battista ([jonathan.battista@enterprisedb.com](mailto:jonathan.battista@enterprisedb.com)), Niccolo Fei ([niccolo.fei@enterprisedb.com](mailto:niccolo.fei@enterprisedb.com)), Gabriele Bartolini ([gabriele.bartolini@enterprisedb.com](mailto:gabriele.bartolini@enterprisedb.com))

The "Provided APIs" section lists five services with "Create instance" buttons:

- Backups:** PostgreSQL backup (physical base backup)
- Cluster Image Catalog:** A cluster-wide catalog of PostgreSQL operand images
- Cluster:** PostgreSQL cluster (primary/standby architecture)
- Image Catalog:** A catalog of PostgreSQL operand images
- Pooler:** Pooler for a Postgres Cluster (with PgBouncer)

The "Description" section notes that the operator has "Main features:".



# Use case

## Create the postgres cluster



# Bootstrap - different ways of creating a cluster

- Create a new cluster from scratch
  - “initdb”: named after the standard “initdb” process in PostgreSQL that initializes an instance
- Create a new cluster from an existing one:
  - Directly (“pg\_basebackup”), using physical streaming replication
  - Directly (logical backup/restore) using pg\_dump and pg\_restore
  - Indirectly (“recovery”), from an object store
    - To the end of the WAL
      - Can be used to start independent replica clusters in continuous recovery
    - Using PITR



# Configure and Install the Postgres cluster

- Prepare for cluster-creation (ensure minio secrets are in place)

```
./04-prepare-cluster
```

- Create a new 3-node cluster by running

```
./05_install_cluster.sh
```

- Check the status of the cluster (using the CNP plugin):

```
./06_show_status.sh
```



# Create table test with 1000 rows

- Once cluster is running ... (minimum the primary) run the script:

```
./07_insert_data.sh
```

Try it for yourself

10:00



# Use case

## Promote & Upgrade the postgres cluster



# Rolling updates

- Update of a deployment with ~zero downtime
  - Standby servers are updated first
  - Then the primary:
    - supervised / unsupervised
    - switchover / restart
- When they are triggered:
  - Security update of Postgres images
  - Minor update of PostgreSQL
  - Configuration changes when restart is required
  - Update of the operator
    - Unless in-place upgrade is enabled



# Check the cluster status

- In terminal **1**: (prepare a terminal for status - and one to run the admin-commands):
  - Run the command  
./06\_show\_status.sh
  - Review the output:
    - check Postgres version: "PostgreSQL Image: quay.io/enterpriseDB/postgresql:**16.2**"
    - check "Continuous Backup status": "**Not configured**"
  - Check the updated cluster configuration - file cluster-example-upgrade.yaml  
less ./yaml/cluster-sample-upgrade.yaml
    - Check Postgres version: "imageName: quay.io/enterpriseDB/postgresql:**16.4**"
    - Check the Backup section



# Run the Promote and Upgrade

- With this step we will:
  - Promote node-2 to become the primary
  - Run the postgres minor update from the version 16.2 to 16.4
  - We will configure the WAL files backup to the S3 storage
- In the web terminal **2**:
  - Check the upgrade status:  
`./06_show_status.sh`
- In the terminal **1**:
  - Run the script:  
`./08_promote.sh`
  - Run the script:  
`./09_upgrade.sh`

## Try it for yourself

05:00



# Use case Backup & Restore



# Backup and Recovery - Part 1

- Continuous physical backup on “backup object stores”
  - Scheduled and on-demand base backups
  - Continuous WAL archiving (including parallel)
  - From primary or a standby
  - Support for recovery window retention policies (e.g. 30 days)
- Recovery means creating a new cluster starting from a “recovery object store”
  - Then pull WAL files (including in parallel) and replay them
  - Full (End of the WAL) or PITR
- Both rely on Barman Cloud technology
  - AWS S3
  - Azure Storage compatible
  - Google Cloud Storage
  - MinIO



# Backup and Recovery - Part 2

- WAL management
  - Object store
- Physical Base backups
  - Object store
  - Kubernetes level backup integration (Velero/OADP, Veem Kasten K10, generic interface)
  - Kubernetes Volume Snapshots



# Kubernetes Volume Snapshot: major advantages

- Transparent support for:
  - Incremental backup and recovery at block level
  - Differential backup and recovery at block level
  - Based on copy on write
- Leverage the storage class to manage the snapshots, including:
  - Data mobility across network (availability zones, Kubernetes clusters, regions)
  - Relay files on a secondary location in a different region, or any subsequent one
  - Encryption
- Enhances Very Large Databases (VLDB) adoption



# Backup & Recovery via Snapshots: some numbers

Let's now talk about some initial benchmarks I have performed on volume snapshots using 3 `r5.4xlarge` nodes on AWS EKS with the `gp3` storage class. I have defined 4 different database size categories (tiny, small, medium, and large), as follows:

Cluster name	Database size	pgbench init scale	PGDATA volume size	WAL volume size	pgbench init duration
<i>tiny</i>	4.5 GB	300	8 GB	1 GB	67s
<i>small</i>	44 GB	3,000	80 GB	10 GB	10m 50s
<i>medium</i>	438 GB	3,0000	800 GB	100 GB	3h 15m 34s
<i>large</i>	4,381 GB	300,000	8,000 GB	200 GB	32h 47m 47s

The table below shows the results of both backup and recovery for each of them.

	Cluster name	1st backup duration	2nd backup duration after 1hr of pgbench	Full recovery time
	<i>tiny</i>	2m 43s		4m 16s
	<i>small</i>	20m 38s		16m 45s
	<i>medium</i>	2h 42m		2h 34m
	<i>large</i>	3h 54m 6s		2m 2s

<https://www.enterprisedb.com/postgresql-disaster-recovery-with-kubernetes-volume-snapshots-using-cloudnativepg>



# Create the full backup

- With this step we will:
  - Create the full backup of the postgres cluster in the MinIO storage:

- In the web terminal 1:

- Run the script:

```
cd /projects/workshop  
./10_backup_cluster.sh
```

- Check the backup status:

```
cd /projects/workshop  
./11_backup_describe.sh
```



# Check Backup in the Openshift Console

- Obtain the OpenShift Console URL from the Slide #56 and open the URL:
- Navigate to:
  - -> Operators
  - -> Installed Operators
  - -> EDB Postgres for Kubernetes Operator
  - -> Go to the Backup section and show the created backup:

The screenshot illustrates the process of creating a PostgreSQL backup using the EDB Postgres for Kubernetes Operator.

**Left Panel (Installed Operators):**

- 1. The "OperatorHub" tab is selected.
- 2. The "EDB Postgres for Kubernetes" operator is listed under "Managed Namespaces".

**Right Panel (Backup Section):**

- 3. The "Backups" tab is selected.
- 4. A new backup named "cluster-user-backup-test" is listed, showing it was completed successfully.

**Bottom Right Corner:**

© EDB 2025 – ALL RIGHTS RESERVED.

# Check Backup in MinIO UI

- Obtain the MinIO URL from the Slide **#56** and open the URL:
- Connect as user **minio** with the password: **edb-workshop**
- The page will appear:

The screenshot shows the MinIO Object Store UI. On the left, there is a sidebar with the following navigation options:

- User
  - Object Browser
  - Access Keys
  - Documentation
- Administrator
  - Buckets
  - Policies
  - Identity

The main area is titled "Object Browser" and shows the contents of a bucket named "cnp". The bucket was created on "Mon, Apr 14 2025 18:54:07 (GMT+2)" and has "PRIVATE" access. It contains "63.6 MiB - 6 Objects". A search bar at the top right says "Start typing to filter objects in the bucket". The object list table has columns for "Name" and "Last Modified". The table shows one item: "cluster-example".

Name	Last Modified
cluster-example	



# Restore the database from the backup

- With this step we will:
  - Create the new cluster cluster-restore
  - Restore the full backup created in the previous step in the new cluster:
- In the terminal 1:
  - Run the restore:  
`./12_restore_cluster.sh`
  - Check the creation status:  
`kubectl get pods -w # after creation stop the execution with <ctrl>+c`
  - Check the table test in the cluster-restore, run the script:  
`oc exec -it cluster-restore-user<X>-1 -- psql -U postgres -c "\d test"`  
`oc exec -it cluster-restore-user1-1 -- psql -U postgres -c "select count(*) from test;"`
  - Delete the cluster-restore-user<x> to avoid resource problems during the workshop:  
`oc delete cluster cluster-restore-user<X>`



# Backup demonstration

- Create the full backup
- Check Backup in the Openshift Console
- Check Backup in MinIO UI
- Restore the database from the backup

Try it for yourself

15:00



# Use case: Failover



# Run failover test

- With this step we will:
  - Delete the primary database of the cluster cluster-example
  - Check the cluster status in the another terminal window
- In the web terminal 1:
  - Run the script:  
`./13_failover.sh`
- In the web terminal **2**:
  - Check the failover cluster status:  
`./06_show_status.sh`

Try it for yourself

05:00



# Use case Scale-out and scale-down



# Scale-out the postgres cluster

- With this step we will:
  - Add the 1 standby to the cluster
- In the web terminal 1:
  - Run the script:  
`./14_scale_out.sh` (using `-replicas=X...` another way would be to update the YAML)
- In the web terminal **2**:
  - Check the cluster status:  
`./06_show_status.sh`



# Scale-down the postgres cluster

- With this step we will:
  - Remove 2 standby pods from the cluster
- In the web terminal 1:
  - Run the script:  
`./15_scale_down.sh`
- In the web terminal **2**:
  - Check the cluster status:  
`./06_show_status.sh`

Try it for yourself

05:00



# Use Case Fencing



# Stop postgres process on the pod

- In the web terminal 1:
  - Run the script:

```
./30_fencing_on.sh
```

- In the web terminal **2**:
  - Check the cluster status:

```
./06_show_status.sh
```



# Start the postgres process on the pod

- In the terminal 1:
  - Run the script:  
`./31_fencing_off.sh`
- In the terminal **2**:
  - Check the cluster status:  
`./06_show_status.sh`

## Try it for yourself

05:00



# Use case Hibernation



# Stop the postgres cluster

- In the terminal 1:
  - Run the script:

```
./32_hibernation_on.sh
```

- In the terminal **2**:
  - Check the cluster status:

```
./06_show_status.sh
```



# Start the postgres cluster

- In the terminal 1:
  - Run the script:

./33\_hibernation\_off.sh

- In the terminal **2**:
  - Check the cluster status:

./06\_show\_status.sh

## Try it for yourself

05:00



# Use case Major version Upgrade



# Delete cluster restore and upgrade cluster

- In the web terminal 1:
  - Run the script:  
./20\_upgrade\_major\_version.sh
  - Check the cluster creation process:  
kubectl get pods -w
  - Check the table test in the cluster-user<X>-17, run the command:  
oc exec -it cluster-user<X>-17-1 -- psql -U postgres -d app -c "\d test"  
oc exec -it cluster-user<X>-17-1 -- psql -U postgres -d app -c "select count(\*) from test;"



What more?  
(some additional features from EDB)

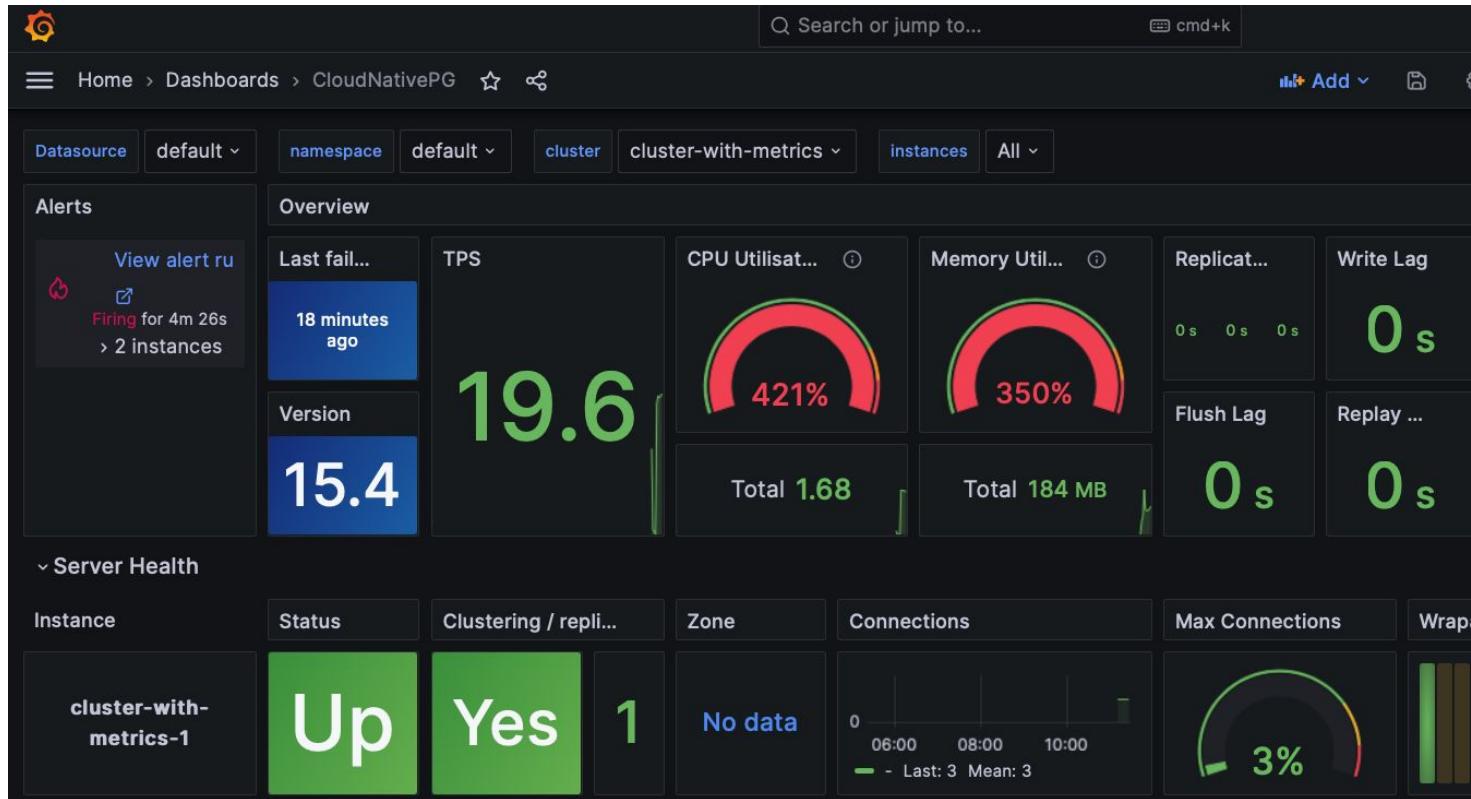


# What we didn't show you today ....

- PgBouncer (Pooler) integration
  - Create a PgBouncer deployment and automatically configure to the cluster.
- Monitoring using Prometheus and Grafana
  - Exporting to OpenMetrics (Prometheus)



# Grafana Dashboard



# Advanced Security



## Password policy management

DBA managed password profiles, compatible with Oracle profiles



## Audit compliance

Track and analyze database activities and user connections



## Virtual private databases

Fine grained access control limits user views



## EDB/SQL protect

SQL firewall, screens queries for common attack profiles



## Data redaction

Protect sensitive information for GDPR, PCI and HIPAA compliance



## Code protection

Protects sensitive IP, algorithms or financial policies

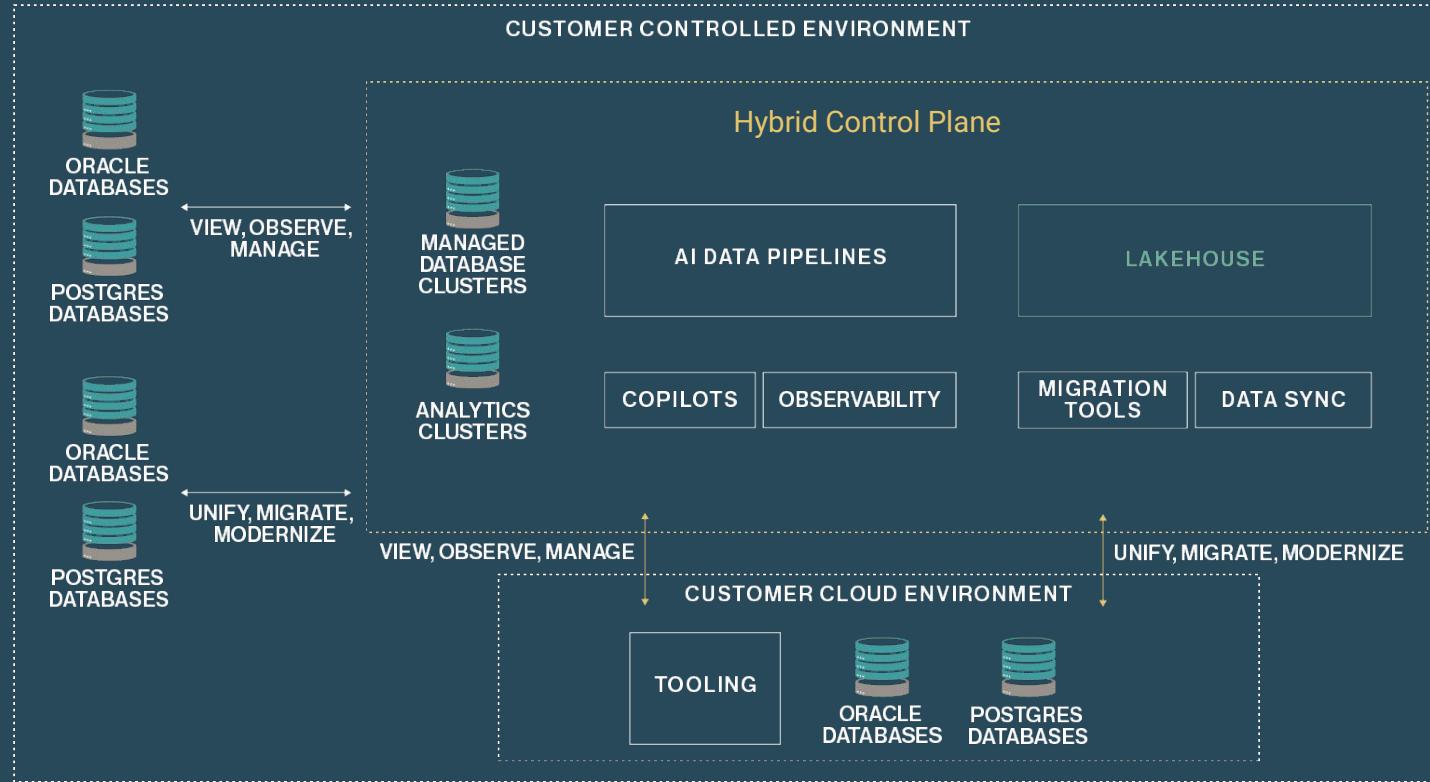


# Transparent Data Encryption (EDB-only features)

- Transparent Data Encryption (TDE) is a feature of EDB Postgres Advanced Server and EDB Postgres Extended Server that prevents unauthorized viewing of data in operating system files on the database server and on backup storage
- Data encryption and decryption is managed by the database and does not require application changes or updated client drivers
- EDB Postgres Advanced Server and EDB Postgres Extended Server provide hooks to key management that is external to the database allowing for simple passphrase encrypt/decrypt or integration with enterprise key management solutions, with initial support for:
  - Amazon AWS Key Management Service (KMS)
  - Google Cloud - Cloud Kay Management Service
  - Microsoft Azure Key Vault
  - HashiCorp Vault (KMIP Secrets Engine and Transit Secrets Engine)
  - Thales CipherTrust Manager
- Data will be unintelligible for unauthorized users if stolen or misplaced

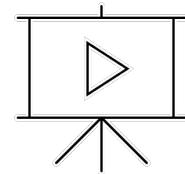


# Hybrid Control Plane at a glance



# Hybrid Control Plane

LIVE DEMO



# Something Big Is Coming

June 17, 2025 | 11 AM CEST

Meet the Future of EDB Postgres® AI



On June 17, we reveal the next stage of EDB Postgres AI and show you how enterprises are using it to build secure, scalable, sovereign AI foundations.

Join us and our partners Red Hat & SuperMicro to explore what's driving that and how EDB Postgres AI helps you act on it.





Share Your Opinion &  
Win a LEGO  
**LEGO Technic 42207  
Ferrari!**

Scan the QR Code & complete the survey  
to let us know your feedback!

*By accepting this prize, you confirm that the prize conforms to your employer's internal rules, policies, and codes of conduct.*

FERRARI SF-24





<https://wheelofnames.com/>



By accepting this prize, you confirm that the prize conforms to your employer's internal rules, policies, and codes of conduct.



Thank you for participating in the  
Postgres on Kubernetes Workshop

Please pick up your certificate :-)

